Ukraine Conflict Similar Tweets

Gabriele Cerizza

Università degli Studi di Milano
gabriele.cerizza@studenti.unimi.it
https://github.com/gabrielecerizza/amd_project

Introduction

In this report we detail our findings in the study of an algorithm capable of identifying similar pairs of documents within massive datasets. The experiments illustrated hereinafter were carried out as part of a project for the Algorithms for Massive Datasets course of Università degli Studi di Milano.

In Section 1 we illustrate the dataset and the adopted pre-processing techniques. In Section 2 we briefly describe the algorithm and its implementation. We also outline the neural network model used to evaluate the performance of the algorithm. In Section 3 we expound on the findings of our experiments. Finally, Section 4 contains our concluding remarks.

1 Dataset

In this section we provide an overview of the dataset (Section 1.1) and of the pre-processing techniques (Section 1.2).

1.1 Description

The dataset employed in our experiments was the "Ukraine Conflict Twitter Dataset" from Kaggle[1], released under the CC BY-NC-SA 4.0 license. We assume the reader to be familiar with the terminology associated with the Twitter platform, including expressions such as "tweet", "retweet", "hashtag", and "handle"[2].

This dataset boasts a total of over 40 million tweets concerning the conflict between Russia and Ukraine, which broke out on the 24th of February 2022. Those tweets were collected daily by monitoring hashtags. The dataset was first published on the 27th of February 2022. In the remainder of this report, we will refer to the version 127 of the dataset, downloaded on the 19th of June 2022.

The dataset comprises 109 compressed CSV files. For each sampled tweet, we can find information concerning the author, the text, the hashtags, the language

---

[1] www.kaggle.com/datasets/bwandowando/ukraine-russian-crisis-twitter-dataset-1-2-m-rows

[2] We refer to https://help.twitter.com/en/resources/glossary for a quick review of the terminology.

and the date of creation of the messages. We used the `text` field to extract each document and the `language` field to select a subset of English documents.

It is worth noting that the naming of the files is not consistent, which hinders attempts to process the tweets chronologically.

1.2   Filtering and Pre-processing

The dataset contains a sizeable number of retweets or tweets that differ only for the inclusion of a number or emoji or handle or punctuation symbol or URL. Tweets consisting only of hashtags or very short messages are likewise abundant. Given the nature of the Twitter platform, this is to be expected. However, such documents do not provide a substantial challenge to models aiming at retrieving similar documents, given the amount of identical, overlapping text. For this reason, we set out to remove as much as possible those documents by way of a first filtering phase. The documents were then further pre-processed before being fed to the models.

Filtering. Filtering was carried out by performing the following operations:

1. we removed URLs;
2. we removed handles;
3. we replaced words with accents with their counterparts without accents;
4. we replaced special UNICODE characters with their ASCII counterparts, for instance "$\mathbb{R}$" with "R";
5. we replaced the ampersand with "and";
6. we converted everything to lowercase;
7. we removed non-ASCII characters, except for cyrillic characters;
8. we removed punctuation symbols;
9. we removed stop words;
10. we normalized the whitespace;
11. we dropped documents shorter than 100 characters.

These steps were performed having two objectives in mind: (i) dropping duplicates, after discarding noisy and irrelevant information; and (ii) reducing the number of different characters that could be found in the tweets, which adversely affects the main algorithm complexity (see further ...).

Pre-processing. Different pre-processing pipelines were applied to the documents given as input to the main algorithm described in Section and to the neural network model described in Section.

Concerning the documents given as input to the main algorithm, the filtering steps described above also directly affected the documents and, thusly, may be considered part of the pre-processing pipeline. However, the filtering steps were also used to drop the documents that, after those operations, turned out to be identical. On the contrary, the pre-processing steps described below were not involved in determining whether two documents should be considered identical. These pre-processing steps were:

1. using the `spacy` library[3], we replaced each token with the corresponding lemma, to increase the match between words like "invasion" and "invaded";
2. we removed punctuation symbols that were introduced by the lemmatization;
3. we normalized the whitespace again.

The documents given as input to the neural network model were not affected by the filtering steps, which, in this case, were used solely to drop duplicates. As a consequence, the only pre-processing steps were the following:

1. we removed URLs;
2. we replaced words with accents with their counterparts without accents;
3. we replaced the ampersand with "and";
4. we replaced special UNICODE characters with their ASCII counterparts, for instance "$\mathbb{R}$" with "R";
5. we removed non-ASCII characters, except for cyrillic characters;
6. we normalized the whitespace.

We decided to perform a light pre-processing for the neural network model in order to leverage its ability to exploit the context of each word, which would have been hampered by removing punctuation and stop words or by lemmatizing the tokens.

Finally, note that, while numbers might be considered noise in some contexts, here we decided to keep them. Indeed, in our context, numbers could be found in dates and military equipment (e.g., the Russian T-72 tank or the M982 Excalibur 155 mm shell) and could help in discriminating the documents.

2   Models

In this section we briefly describe the algorithm used to find similar documents (Section 2.1) and the neural network model we used as a baseline to compare performances (Section 2.2). We also propose an example to motivate the comparison between the two models (Section 2.3).

2.1   Locality-sensitive Hashing Algorithm

In order to find the pairs of similar documents, we employed the algorithm described in [3], which can scale up to massive datasets. This algorithm converts each document to a set of k-grams (shingles), builds a characteristic matrix and then a signature matrix by using hash functions, applies a locality-sensitive hashing (LSH) technique to find candidate similar pairs and, finally, checks the candidate pairs against the signature matrix to discard false positive pairs. Optionally, one could check the candidate pairs also against the characteristic matrix, if available.

---

[3] https://spacy.io/

Hash functions. Concerning the implementation details, we first discuss the various hash functions exploited by the algorithm. To avoid keeping the shingles in main memory, the algorithm hashes each shingle to a bucket, in the form of an integer. The total number of buckets associated to the hash function has a significant impact on the execution time. We generated different hash functions by truncating the output of the SHA-256 hash function at varying lengths of bits. In this way, we could experiment with different numbers of buckets.

The algorithm requires also hash functions that could map a row index to another row index. These hash functions are used to efficiently perform permutations of the rows. Since the number of these hash functions is a hyper-parameter set by the user, we needed a way to generate an arbitrary number of different hash functions. To this end, we followed the approach described in [4], generating hash functions of the form

$$h(x) = (ax + b) \bmod c \ ,$$

where x is a row number, a and b are random numbers smaller than the maximum row number and c is a prime number higher than the maximum row number. Note that a and b must be unique for a given signature matrix.

Inside the LSH technique, to check whether two columns were equal, we hashed the string composed by the sequence of integers inside each column using the default `hash` function provided by Python. This allowed us compress a potentially huge array of numbers into a single number, thus saving main memory space.

Characteristic matrix. In order to build the signature matrix, the algorithm also requires a characteristic matrix, where each column represents a document and each entry indicates whether or not a given shingle is contained in a given document. We store the characteristic matrix as a dictionary that maps each document index to its set of shingles. In this way, we do not waste memory to store the zeroes.

LSH. Concerning the LSH technique, we found the number of bands for the signature matrix and the correlated number of rows in each band by using `scipy` to solve the system of equations

$$\begin{cases} t = \left(\frac{1}{b}\right)^{\frac{1}{r}} \\ b \cdot r = n \end{cases} , \tag{1}$$

where t is a threshold on the required similarity between two documents to mark them as a candidate pair, b is the number of bands, r is the number of rows in each band and n is the total number of rows of the signature matrix.

As noted before, we checked whether the columns were equal by hashing their representations as strings.

Scalability bottlenecks. The implementation of the algorithm features two possible scalability bottlenecks. The first one concerns the characteristic matrix. Using 3-grams, 4-grams, 5-grams and 6-grams, each document contained on average 120-130 shingles. If we represented each shingle with a 32-bit integer, processing 10 million documents would require roughly 5 gigabytes of main memory ($\frac{1e7 \cdot 130 \cdot 32}{8e9} = 5.2$). Keeping the whole dataset (40 million documents) in main memory might prove challenging or even impossible. A solution could be to store the documents representations as sets of shingles in mass memory and read them one at a time when building the signature matrix. However, that would slow down an already time consuming operation.

The second bottleneck concerns the number of similar pairs. The lower the value for the threshold, the higher the number of similar pairs identified by the algorithm. When processing millions of documents, even a moderate threshold might saturate the main memory. The only possible solution would be to store the candidate pairs identified by the LSH algorithm on mass memory. This would happen within each band and, therefore, we would not be able to have unique pairs: the same pair might be identified and stored on mass memory in more than one band. This will result in redundant processing.

2.2   Neural Network (Transformer)

The dataset did not provide a ground truth against which we could compare our results. Therefore, we decided to use a state-of-the-art model in the field of natural language processing (NLP) to serve as a baseline in determining how similar two documents were. To this end, we opted to use MPNet [6], which is a neural network based on the revolutionary BERT Transformer [1, 7]. A characteristic of the Transformer architecture is the computation of attention weights, which boils down to learning the relative importance of each token inside a sentence.

A Transformer used to obtain word embeddings is similar to word2vec [5], with one important difference. The difference is that word2vec produces static embeddings, which means that a given word will be mapped to the same vector regardless of its context. On the contrary, the Transformer produces contextualized embeddings, so that we will obtain a different vector for each different context (sentence) in which a given word appears. To give an example, the Transformer would generate different embeddings for the word "mouse" used with the meaning of animal or with the meaning of computer device, whereas word2vec would generate the same embedding.

We used the Transformer to obtain an embedding for each document. First, we extracted a word embedding for each word in a given document by using the last hidden layer of the model, which captures semantic features [2]. Afterwards, we computed the mean of the word embeddings of the document.

2.3 Motivation

We motivate our decision to compare the LSH algorithm with a Transformer by showing an example. Consider the following three documents.

1. I went to the bank to withdraw money and compensate the plaintiff for their losses.
2. The Russians are withdrawing from the banks of the Dnipro River after suffering heavy losses.
3. After the judgement, I had to use my debit card and pay the suer for the damages.

We can see that the first and the third document are semantically related, despite the fact that they do not share many words in common. On the contrary, the first and the second document are not semantically related, but share words like "bank", "withdraw" and "losses".

   We apply a light pre-processing on the documents, mainly consisting in removing punctuation and performing lemmatization, we extract 3-grams from the texts and then we compute the Jaccard similarity between each pair of documents. Recall that the LSH algorithm produces an estimation of the Jaccard similarity. After that, we compute the cosine similarity between the sentence embeddings extracted from the Transformer, without any pre-processing.

   The results are shown in Table 1. As expected, the Jaccard similarity identifies the first and second document as the most similar, whereas the Transformer correctly identifies the first and the third document as the most similar. This suggests that we could use the Transformer to check whether the pairs identified as similar by the LSH algorithm are indeed similar.

Table 1. Similarity between each pair out of the three documents described in Section 2.3, computed by Jaccard similarity on 3-grams and by the Transformer.

| Method	Similarity		
	Pair (1,2)	Pair (1,3)	Pair (2,3)
Jaccard Similarity	0.185	0.000	0.023
Transformer	0.192	0.648	0.131

3 Experiments

In this section we disclose and discuss our experiments on the LSH algorithm. These experiments concerned the growth of shingles and number of characters as we increased the number of documents (Section 3.1); the effect of the number of buckets on the hash functions used to map shingles to integers (Section 3.2); the effect of the threshold on the similarity between pairs (Section 3.3); the effect

of the number of hash functions used to permute the rows of the characteristic matrix (Section 3.4); a final comparison between the LSH algorithm and the Transformer on 100,000 tweets (Section 3.5).

All the experiments were carried out using the files covering the tweets from the 1st of April 2022 to the 7th of April 2022. We selected only the tweets in the English language. We refer to Section 1.2 for the filtering and pre-processing techniques adopted. The analyzed files contained around 2 million documents. After the filtering, only 221,615 documents remained. This confirms that the dataset is filled with almost identical or very short tweets.

3.1   Shingles and Characters Growth

Fig. 1. Growth in the number of characters as the number of documents increases.

In this first experiment we wanted to gauge how the number of shingles (or k-grams) and different characters within the whole corpus varied as the number of documents increased.

In Figure 1 we can see how the number of characters grows as the number of documents increases. In particular, we can see how from 100,000 documents onwards the number of characters is stable at around 65 characters (68, to be precise). This is useful to get a hint on the appropriate k-gram with respect to the number of buckets of the hash function that maps k-grams to integers. For
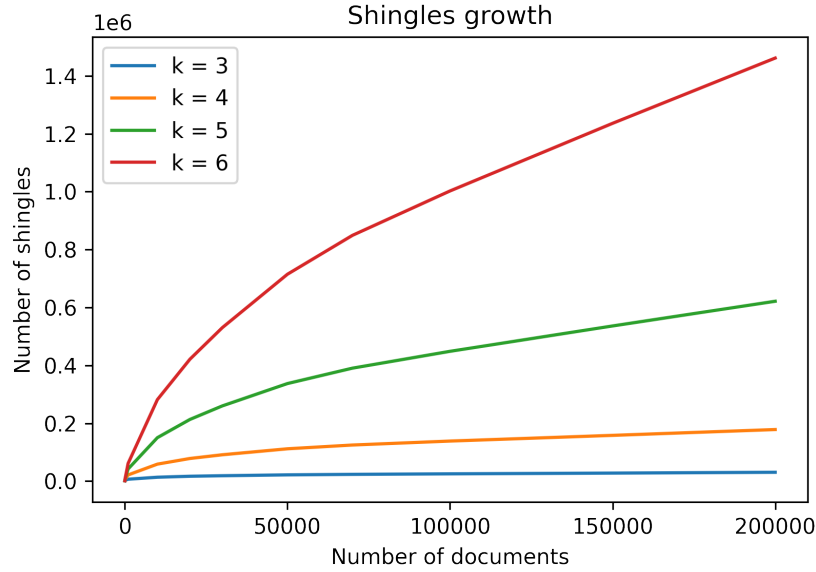
Fig. 2. Growth in the number of shingles as the number of documents increases. The growth is shown for shingles obtained with different k-grams.

instance, given that $68^3 = 314,432$, we can infer that, if we used 3-grams, we would need a hash function with approximately 300,000 buckets.

Working with characters we can get an estimate on the number of buckets that we would need. However, it is unlikely that all the possible shingles will appear in practice. To this end, we also explored how the actual number of different shingles grows as the number of documents increases.

In Figure 2 we can see how the number of different shingles grows as the number of documents increases. It is not surprising to see that, with k-grams having a higher value of k, the number of shingles grows more steeply. We can also see, however, that this growth is approximately logarithmic in shape and stabilizes quickly for low values of k. Indeed, for k equal to 3 and 4, most of the shingles have already been found at 50,000 documents. To be precise, at 50,000 documents we have 21,580 shingles with 3-grams and 111,643 shingles with 4-grams. For 3-grams, a hash function with buckets of at least 15 bits would be enough ($2^{15} = 32,768$ buckets). For 4-grams, we would need at least 17 bits to have a fair chance at avoiding collisions ($2^{17} = 131,072$ buckets).

3.2   Number of Buckets

In this experiment we wanted to see how increasing the number of bits, and thus the number of buckets, of the hash function that maps k-grams to integers affected various metrics. The experiment was carried out with a fixed threshold

of 0.1, 100 hash functions for the permutations of the rows and a corpus of 100 documents.

The first metric we examined was precision, defined as

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}},$$

where TP is the number of true positives and FP is the number of false positives. TP and FP are obtained by computing the actual Jaccard similarity on the pairs of documents identified as similar by the algorithm. This was done using the characteristic matrix. Afterwards, we checked if the Jaccard similarity was indeed higher than the threshold. If the Jaccard similarity was higher than the threshold, we considered the pair a true positive, otherwise we considered the pair a false positive.

The second metric we examined was the mean absolute error (MAE) between the similarity estimated by the algorithm and the actual Jaccard similarity. MAE is defined as

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^{N} |y_i - x_i|,$$

where x and y are the predictions and ground truth values.

As third metric, we examined how the number of bits affected the time of execution.

The results are shown in Tables 2, 3 and 4 for 3-grams, 4-grams and 5-grams, respectively. We can formulate the following observations.

- Using a higher number of bits did not yield a better result. Indeed, we can see that with 3-grams 12 bits yielded the best result, while with 4-grams 12 bits yielded a better result than 22 bits. Although the corpus was limited to 100 documents, this might still suggest that collisions have only a mild impact on the performance.
- Across all the different numbers of bits, precision decreased as the the value of k in k-grams increased.
- No significant pattern was discernible for MAE, either considering the number of bits or the different k-grams.
- The higher the value of k in k-grams, the lower the number of pairs estimated as similar.
- Increasing the number of bits raised the execution time. With 22 bits, processing 100 documents took roughly 29 minutes. This might make working with massive datasets prohibitive from the perspective of execution time. Note also that, for low values of bits, most of the time was spent reading the files.

3.3 Threshold

In this experiment we wanted to see how the value of the threshold affected the predictions, considering precision and MAE as metrics. The experiment was

Table 2. Effects of the increase in the number of bits for the hash function that maps k-grams to integers when k is equal to 3. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash Bits	True Positive	False Positive	Precision	MAE	Time Delta
12	1289	560	**0.697**	0.021	0:02:06
14	945	738	0.561	0.023	0:02:11
16	705	328	0.682	**0.018**	0:02:28
18	722	374	0.659	0.020	0:03:45
19	767	500	0.605	0.019	0:05:27
20	779	413	0.654	0.019	0:08:50
22	713	457	0.609	0.020	0:28:55

Table 3. Effects of the increase in the number of bits for the hash function that maps k-grams to integers when k is equal to 4. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash Bits	True Positive	False Positive	Precision	MAE	Time Delta
12	131	256	0.339	0.024	0:02:09
14	65	130	0.333	0.022	0:02:12
16	66	140	0.320	0.024	0:02:29
18	52	121	0.301	**0.018**	0:03:41
19	36	41	**0.468**	**0.018**	0:05:26
20	46	56	0.451	**0.018**	0:08:41
22	69	245	0.220	0.021	0:29:08

Table 4. Effects of the increase in the number of bits for the hash function that maps k-grams to integers when k is equal to 5. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash Bits	True Positive	False Positive	Precision	MAE	Time Delta
12	54	187	0.224	0.023	0:02:09
14	32	165	0.162	0.021	0:02:12
16	29	26	**0.527**	0.021	0:02:31
18	23	82	0.219	**0.019**	0:03:47
19	25	151	0.142	0.034	0:05:28
20	31	76	0.290	0.021	0:08:44
22	18	24	0.429	0.021	0:29:18

carried out with a 16 bits hash function for the mapping between shingles and integers, 100 hash functions for the permutations of the rows and a corpus of 100 documents. Note that, as the 100 hash functions are generated randomly, the results might be slightly different from the ones obtained in the previous experiment, which was carried out with a threshold of 0.1.

The results are shown in Tables 5, 6 and 7 for 3-grams, 4-grams and 5-grams, respectively. We can formulate the following observations.

– Above a certain threshold, the algorithm reached a perfect precision for all the k-grams. Below that threshold, however, increasing the threshold did not always increase the precision.
– No clear pattern was discernible for MAE.
– Increasing the threshold decreased the number of estimated similar pairs, as expected. However, this effect was stronger for k-grams with higher values of k.

Table 5. Effects of the increase in the threshold value when working with 3-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Threshold	True Positive	False Positive	Precision	MAE
0.05	3740	263	0.934	**0.019**
0.10	924	643	0.590	0.021
0.15	79	449	0.150	0.036
0.20	4	85	0.045	0.035
0.25	2	3	0.400	0.052
0.30	1	1	0.500	0.020
0.50	1	0	**1.000**	0.060

Table 6. Effects of the increase in the threshold value when working with 4-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Threshold	True Positive	False Positive	Precision	MAE
0.05	1255	838	0.600	0.019
0.10	76	189	0.287	0.021
0.15	7	10	0.412	0.021
0.20	3	2	0.600	**0.015**
0.25	1	1	0.500	0.102
0.30	1	0	**1.000**	0.028
0.50	1	0	**1.000**	0.058

Table 7. Effects of the increase in the threshold value when working with 5-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Threshold	True Positive	False Positive	Precision	MAE
0.05	423	540	0.439	0.018
0.10	19	22	0.463	0.017
0.15	3	3	0.500	0.012
0.20	1	0	**1.000**	0.111
0.25	1	0	**1.000**	**0.011**
0.30	1	0	**1.000**	0.041
0.50	1	0	**1.000**	0.021

3.4   Number of Hash Functions

In this experiment we wanted to verify how the number of hash functions used to perform permutations affected precision, MAE and execution time. The experiment was carried out with a 16 bits hash function for the mapping between shingles and integers, a threshold value of 0.1 and a corpus of 100 documents.

The results are shown in Tables 8, 9 and 10 for 3-grams, 4-grams and 5-grams, respectively. We can formulate the following observations.

- Increasing the number of hash functions decreased the number of pairs estimated as similar and increased the precision. The effect is similar to what we observed in the threshold experiment, but the meaning is entirely different. Since the threshold here is fixed, when we predict a lower number of true positive pairs we are effectively increasing the number of false negatives. For this reason, 200 hash functions seem to strike a good compromise between precision and the preservation of true positives.
- No clear pattern is discernible for MAE.
- The increase in execution time, when working with a higher number of hash functions, is negligible.

Table 8. Effects of the increase in the number of hash functions used to perform permutations when working with 3-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash No.	True Positive	False Positive	Precision	MAE	Time Delta
20	868	1338	0.393	0.052	0:02:25
100	1183	1409	0.456	0.033	0:02:27
200	812	370	0.687	0.015	0:02:33
300	70	5	**0.933**	**0.014**	0:03:48
500	0	0	0.000	0.000	0:02:47

Table 9. Effects of the increase in the number of hash functions used to perform permutations when working with 4-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash No.	True Positive	False Positive	Precision	MAE	Time Delta
20	32	317	0.092	0.049	0:02:27
100	55	94	0.369	0.022	0:02:30
200	52	32	**0.619**	**0.016**	0:02:35
300	6	3	0.667	0.018	0:02:47
500	0	0	0.000	0.000	0:02:50

Table 10. Effects of the increase in the number of hash functions used to perform permutations when working with 5-grams. We highlighted in bold the best results for precision and mean absolute error (MAE).

Hash No.	True Positive	False Positive	Precision	MAE	Time Delta
20	25	529	0.045	0.044	0:02:30
100	22	48	0.314	**0.018**	0:02:31
200	24	27	0.471	**0.018**	0:02:43
300	1	0	**1.000**	0.028	0:02:54
500	1	0	**1.000**	0.027	0:02:55

3.5   Similarity between 100,000 Tweets

For this last experiment, we compared the results of three different configurations of algorithm hyper-parameters against the Transformer. The Transformer would not be able to compute the pairwise similarity between each pair of documents in a massive dataset. That would require an unreasonable amount of time. As a consequence, we used the Transformer to compute the similarity only between the pairs selected by the LSH algorithm. After that, we compared the rank-order correlation between the similarities computed by the two models using Kendall's Tau and Spearman's Rho.

Results were still acceptable and required much less time

4   Conclusions

This work investigated the performance of the LSH algorithm described in Section 2.1 across different configurations of hyper-parameters. The algorithm proved to be effective at processing massive amounts of data and at properly

estimating the Jaccard similarity between documents. The main pitfall of the algorithm is that Jaccard similarity on k-grams, as shown in Section 2.3, might not capture semantic similarity between documents. Given the fact that the algorithm is built on top of a Boolean matrix, this issue may prove difficult to address. The reason is that we cannot exploit state-of-the-art word embeddings, unless we find a way to suitably encode them as sets.

Another critical aspect of the algorithm was the execution time, as mentioned in Section 3.2. Building the signature matrix is an expensive operation, especially when using hash functions with a high number of buckets. In [3] there are suggestions on how to address this issue, at the cost of a coarser approximation of the Jaccard similarity.

As a further improvement, we could use a non-cryptographic hash function, instead of SHA-256, to generate hash functions with a variable number of buckets. That should speed up the execution.

We could improve the algorithm by appending argument type constraints to the templates.

MAE ON TP AND FP, OR ALSO TN AND FN

INCONSISTENT RESULTS FOR THRESHOLD AND NUMBER OF BITS (3-GRAMS, thresh 0.1). Different hash though. the probability that any prefixed shingle will be contained in my document will be small $= 280/68^3$ SET SEED WHEN CREATING HASH

ADD ACCURACY?

After feeding the documents to

Example of similar tweets Non-cryptographic function

and which has been found to work well Link a zip with checkpoints in the github To achieve this goal To this end

, so we removed all URLs beforehand. We also removed URLs from tweets, since that differed only in hashtags or in placed within the messages were not filtered, since .

might skew

USE COMMA NOT PERIOD for NUMBERS

each of which organized in fields providing information about the author each of which

foregoing

Not actually deduped.

Optimizer to find b and r Hash in LSH Last layer has most meaning TFIDF? Need to traverse the whole dataset to compute, taxing operation Two semantically different pipelines: to filter almost identical tweets, and the other. Punctuation and non-ascii removed later, because they might help the lemmatizer CHANGED PIPELINE (some differed for only some whitespace or for a capital letter (mint and MINT))

Problem with low threshold, ends RAM with combinations As such, we decided to experiment with The algorithm used dictionary to store shingles

List of documents, each with its own correlation Check TP on final

Being confronted with the problem of finding hash functions that could map to different numbers of buckets,

filled to the brim of duplicates no clear takeaways more hashes more FN?

k=3 with t0.4 over 1M pairs

Strip again punctuation because some of the lemmas had punctuation symbols like _ in them.

Retweets and Quote tweets. We assume the reader to be familiar with Twitter.

Check if quote or retweet with quote and retweet status (https://developer.twitter.com/en/docs/twitter-api/v1/data-dictionary/object-model/tweet)

Filtering for followers number doesn't seem useful. Many people retweet the smae things. maybe they made mistakes.

fast spike increase in shingles at low documents

used spacy

Cannot use spacy on dataframe, takes too much time

crash with t.02, too many pairs with t=0.4 6346167 candidate pairs, actual 342671

We do not drop duplicates across different files, as that would require to load in main memory millions of documents. Plenty of retweets. Plenty of very short tweets that are identical after preprocessing /lemmatization. But filtering on string length isn't suitable due to the instrinsic short nature of tweets. Sometimes only changed a handle or more, but removing handles might remove important parts (many quote zelensky with handle for instance).

Removing handles might be somehow questionable... Si ma poi se tanto retweettano in altri file compare comunque.

Dai, remove handle e testi più lunghi di tot... ONLY EN language Should have kept the retweet flag

CHECK TP ON ACTUAL OUTPUT, AFTER THRESH CHECK

Generate the combinations manually and add one at a time, to avoid ram problem?

References

1. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: BERT: Pre-training of deep bidirectional transformers for language understanding. In: Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers). pp. 4171–4186. Association for Computational Linguistics, Minneapolis, Minnesota (Jun 2019). https://doi.org/10.18653/v1/N19-1423, https://aclanthology.org/N19-1423
2. Laicher, S., Kurtyigit, S., Schlechtweg, D., Kuhn, J., Schulte im Walde, S.: Explaining and improving BERT performance on lexical semantic change detection. In: Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Student Research Workshop. pp. 192–202. Association for Computational Linguistics, Online (Apr 2021). https://doi.org/10.18653/v1/2021.eacl-srw.25, https://aclanthology.org/2021.eacl-srw.25

3. Leskovec, J., Rajaraman, A., Ullman, J.D.: Mining of Massive Datasets. Cambridge University Press, 3 edn. (2020). https://doi.org/10.1017/9781108684163
4. Liu, E.: Text Similarity. http://ethen8181.github.io/machine-learning/clustering_old/text_similarity/text_similarity.html (2015), [Online; accessed 23-June-2022]
5. Mikolov, T., Sutskever, I., Chen, K., Corrado, G.S., Dean, J.: Distributed representations of words and phrases and their compositionality. In: Burges, C.J.C., Bottou, L., Welling, M., Ghahramani, Z., Weinberger, K.Q. (eds.) Advances in Neural Information Processing Systems. vol. 26. Curran Associates, Inc. (2013)
6. Song, K., Tan, X., Qin, T., Lu, J., Liu, T.Y.: Mpnet: Masked and permuted pre-training for language understanding (2020). https://doi.org/10.48550/ARXIV.2004.09297, https://arxiv.org/abs/2004.09297
7. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., Polosukhin, I.: Attention is all you need (2017). https://doi.org/10.48550/ARXIV.1706.03762, https://arxiv.org/abs/1706.03762